

L05.01

Hashing II

50.004 Introduction to Algorithms
Ioannis Panageas (ioannis@sutd.edu.sg)
ISTD, SUTD
CLRS Ch 11.3-11.4
Slides by A.Binder and based on Dr. Simon LUI

Recap last lecture ...

- Chained hash:
 - an array s.t. each entry is a linked list
 - a hash function that maps keys onto array indices
 - Map n keys from a very large key space onto hash table of size m
- may have hash collisions
 - Hash operations: insert, delete as $O(1)$ worst case
 - Search $O(n)$ worst, $O(1+\alpha)$ average case (simple uniform hashing assumption)
 - When combined with table doubling: $O(1)$ average case
 - Amortized costs for table doubling+ insertions: also $O(1)$
- Good hash functions: close to **simple uniform hashing assumption** each key has equal chance to end up in any of the bins ... ensures on average equal load across the whole table

Today

- Open addressing – hashing without linked lists
 - Different probing strategies
- Cuckoo hashing $O(1)$ search worst case
- Other uses of hashes / cryptographic hashes
 - File tampering
 - Digital message signing with hashed messages
- Universal hashing


Worse case discussion

Assume $|U|=N$, hash table size = l , $N > l$

Lemma: Then for any hash function h there exist **at least N/l keys** that hash on the same position (collide)

perfectly uniform

distribution of N keys
over the m slots



For **any** hash function we can find $\geq N/l$ keys that collide

- An alternative to deal with collisions than doubly linked lists?

Open addressing

- Do away with doubly linked lists, one array entry = one element
- Pay a price:
 - can fill table with at most $n-1$ elements, then must do table doubling ...
 - Need to search for free slots if $A[h(x.key)]$ is used “Probing”
 - Replace hash function $h(k)$ by a hash function $h(k,i)$ with a parameter i that allows to search for the next slot

Example – linear probing

- k – key to be inserted
- $g(k) = k \bmod 10$
- $h(k,i) = (g(k)+i) \bmod 10$ i – index for probing, start with $h(k,i=0)$

	<i>After insert 89</i>	<i>After insert 18</i>	<i>After insert 49</i>	<i>After insert 58</i>	<i>After insert 9</i>
0			49	49	49
1				58	58
2					9
3					
4					
5					
6					
7					
8		18	18	18	18
9	89	89	89	89	89

Open addressing algorithms

- Linear probing

$$h(k, i) = (h'(k) + i) \bmod m$$

- Quadratic probing

$$h(k, i) = (h'(k) + c_1 i + c_2 i^2) \bmod m$$

- Double hashing

$$h(k, i) = (h_1(k) + i h_2(k)) \bmod m$$

$$h_1(k) = k \bmod 13$$

$$h_2(k) = 1 + k \bmod 11$$

$$k = 14 : h_1(14) = 1, h_2(14) = 4$$

We probe positions 1, 1+4, 1+2x4, 1+3x4, ...

Example – linear probing

- $h(k,i) = (h'(k)+i) \bmod 11$
- $H'(k) = k \bmod 11$

$$h(28)=28\%11=6$$

$$h(47)=47\%11=3$$

$$h(20)=20\%11=9$$

$$h(36)=36\%11=3$$

$$h(43)=43\%11=10$$

$$h(23)=23\%11=1$$

$$h(25)=25\%11=3$$

$$h(54)=54\%11=10 \text{ insert } 28,47,20$$

0	
1	
2	
3	47
4	
5	
6	28
7	
8	
9	20
10	

0	
1	
2	
3	47
4	36
5	
6	28
7	
8	
9	20
10	

insert 36

0	
1	23
2	
3	47
4	36
5	
6	28
7	
8	
9	20
10	43


insert 43,23

0	54
1	23
2	
3	47
4	36
5	25
6	28
7	
8	
9	20
10	43

insert 25,54

Example – quadratic probing

- $h(k,i) = (h'(k) + c_1i + c_2i^2) \bmod 7$ $c_1=0$, $c_2=1$
- $H'(k) = k \bmod 11$



Questions

A. Linear probing

Consider the hash table in the picture with some keys already inserted. Where will you insert $k=17$ when $h'(k) = k \bmod 13$?

1. location 2
2. location 6
3. location 10

B. Quadratic probing

Where will you insert $k=17$ when $h'(k) = k \bmod 13$, $c1 = 2$, $c2 = 1$?

1. location 2
2. location 7
3. location 12

C. Double hashing

Where will you insert $k=14$ when $h1(k) = k \bmod 13$, $h2(k) = 1 + (k \bmod 11)$?

1. location 2
2. location 9
3. location 10

0	
1	79
2	
3	
4	69
5	98
6	
7	72
8	
9	
10	
11	50
12	

Open Addressing performance

- **Uniform hashing assumption:** each key is equally likely to have any of $m!$ permutations of $\{0, \dots, m-1\}$ as probe sequence
- With above: Insertion: $1/(1 - \alpha)$ probe steps on average
- Search:
 - $1/(1 - \alpha)$ probe steps on average for unsuccessful search
 - $1/\alpha \log(1/(1 - \alpha))$ probe steps on average for successful search

Open addressing vs chaining

- Open addressing:
 - Better cache performance (no pointers to off regions needed when objects are “small”, e.g. integers, floats)
- Chaining:
 - Less sensitive to hash function choices
 - When keys are clustering in parts of the table, then linear/quadratic probing will have many steps
 - Less sensitive to high load factors
 - In practice open addressing needs alpha to be kept small, rule of thumb: like 50-70%, otherwise $1/(1-\alpha)$ becomes a nightmare

Cuckoo Hashing

- Uses 2 hash functions $h_1(k), h_2(k)$, where $h_1(k) \neq h_2(k)$
- Key k stored either in $T[h_1(k)]$ or in $T[h_2(k)]$
- Lookup: Just look at at most 2 places! $O(1)$

- Insertion:

If $T[h_1(k)]$ empty, store the key there


else if $T[h_2(k)]$ empty, store the key there

If both full, store key in $T[h_1(k)]$ and move key that was there to its other location (bumping out the key that might be there, etc)

- If $\alpha < 1$ insertion succeeds with high probability
- If insertion loops: **rehash the entire table** (or double table size)
- Insertion takes constant time on average

Example of cuckoo hashing

insert x
 $h_1(x) = 0$
 $h_2(x) = 4$



$$h_1(A) = 0, h_2(A) = 3$$

$$h_1(B) = 3, h_2(B) = 8$$

x bumps A , A bumps B , B finds an empty location

Bumping H creates a cycle \rightarrow rehash everything

Cuckoo Hashing

- Search is $O(1)$ worst-case, not average
- Insertion is $O(n)$ worst-case (average performance is better)

Other hashing use cases

- File modification check
 - Was your data x tampered that is stored somewhere?
 - Compute hash before you upload,
 - Check by rehashing
 - Problem: when an attacker succeeds to fool you ?

Other hashing use cases

- File modification check
 - Was your data x tampered that is stored somewhere?
 - Compute hash before you upload,
 - Check by rehashing
 - Problem: when an attacker succeeds to fool you ?
 - if he finds an x' such that $h(x)=h(x')$ for the given x

- Digital signatures

- A has public key PK_A , private key SK_A . A can sign a message M by private key to obtain a signature s:
- For large messages a hash $h(M)$ instead of M is signed.
 $s = \text{sign}(h(M), SK_A)$.
- recipient can verify that M was signed by A.
 - B runs a function $\text{verify}(h(M), s, PK_A)$
- Problem: attacker wants to pretend that Alice signed document D2, that the attacker owns. what can he try to do?
- Side info: attacker does not want to show D2 to Alice, no way to let alice sign D2 directly.

- Digital signatures

- A has public key PK_A , private key SK_A . A can sign a message M by private key to obtain a signature s :
- For large messages a hash $h(M)$ instead of M is signed.
$$s = \text{sign}(h(M), SK_A) .$$
- recipient can verify that M was signed by A.
 - B runs a function $\text{verify}(h(M), s, PK_A)$
- Problem: attacker can try to find a document $D1$ such that $h(D1)=h(D2)$, then ask Alice to sign $D1$ to obtain s
- Then reuse s with $D2$

- Commitment check:
 - I want to assure somebody that I committed a sum x of SGD for some project, and that I did not change that sum afterwards on the bank account.
 - I do not want to disclose the sum.
 - Give the person the right to ask the bank to see a hash $z=h(x)$ of the account balance instead
 - Problem: don't want that my ominous partner can reverse the hash, e.g. find that x that created the hash value z .

- Commitment check:
 - I want to assure somebody that I committed a sum x of SGD for some project, and that I did not change that sum afterwards on the bank account.
 - Problem: don't want that partner can reverse the hash, e.g. find the sum x that created the hash value.
 - One solution: I do not hash x but I hash $x+c$, where c is a large random number that I will remember
 c is a so called salt → salted hashing (off lecture)

Cryptographic hashing functions

- What properties hash functions are desirable for such applications?
- **One-way**: given a hash z , it should be infeasible to find the x that created this hash: $h(x)=z$
- **Collision-resistance**: infeasible to find any pair x, x' such that $h(x)=h(x')$
- **Target-collision-resistance**:
 - Given some x it is infeasible to find an x' such that $h(x')=h(x)$
- Desired property: hash maps 2 close keys x, x' to very different locations

Universal hashing

- Problem: If hash function is known, then a malicious creator of keys can force $O(n)$ insertion behaviour
- Solution: choose a function h at random from a function class $\mathcal{H} = \{h_1, \dots, h_r\}$
- What properties should that class have?

Universal hashing

- \mathcal{H} is universal collection of hash functions if for every fixed pair of keys $k_1 \neq k_2$ the number functions h causing a collision is bounded as:

$$|\{ h \in \mathcal{H}: h(k_1) = h(k_2) \}| \leq \frac{|\mathcal{H}|}{m}$$

- Theorem: for h drawn randomly from a uniform distribution over a universal collection of hash functions we have $O(1+\alpha)$ insertion time

Importance of that theorem?

- Theorem: for h drawn randomly from a uniform distribution over a universal collection of hash functions we have $O(1+\alpha)$ insertion time
- What does that mean for our search = $O(1)$ result obtained by table doubling and amortization analysis?

Importance of that theorem?

- Theorem: for h drawn randomly from a uniform distribution over a universal collection of hash functions we have $O(1+\alpha)$ search time
- What does that mean for our search = $O(1)$ result obtained by table doubling and amortization analysis?
- Can **exchange** in last lecture **simple uniform hashing** assumption with **universal collection of hash functions** !!! **$O(1)$ still holds!**


Extra (out of syllabus): example of universal hash function

Universal hash functions

$h(k) = ((\textcolor{red}{a}k + \textcolor{red}{b}) \bmod p) \bmod m$ $m = \text{size of hash table} = \text{arbitrary}$

$p = \text{large prime}$ s.t. all keys are in range $\{0, \dots, p-1\}$

$a \in \{1, \dots, p-1\}$, $b \in \{0, \dots, p-1\}$, are chosen uniformly at random



We can prove: $\text{Prob}_{a,b} \{h_{a,b}(k) = h_{a,b}(l)\} \leq \frac{1}{m}, \forall k, l$

random choice of a, b for each experiment

\Rightarrow less than $\frac{|H|}{m}$ hash functions are "bad" for a pair k, l .
 \rightarrow definition of universality!

(proof, out of curriculum)

- $r = (ak_1 + b) \bmod p, s = (ak_2 + b) \bmod p$
- Claim $r \neq s$

By assumption: $k_1, k_2 \leq p-1$, so $k_1 - k_2$ in $[-(p-1), p-1]$, so $k_1 - k_2 \bmod p \neq 0$, then. Also $a \bmod p \neq 0$, therefore (p must be prime for that)

$a(k_1 - k_2) \bmod p \neq 0$... and this is $r - s \bmod p$.

(a, b) with $a \neq 0 \leftrightarrow r, s$ with $r \neq s$ ($p(p-1)$ elements) because $a = (r - s)(k_1 - k_2)^{-1}$ (this inverse exists in \mathbb{Z}/\mathbb{Z}_p)

... its in CLRS, you need to know about multiplication of equivalence classes in Rings \mathbb{Z}/\mathbb{Z}_p and fields (if p is a prime number, then the ring \mathbb{Z}/\mathbb{Z}_p is a field, i.e. every class has an inverse A bit mathy+technical stuff)

Example of usage

- Built into most modern programming languages (Python, Perl, Java, C++,...)
- Example:
 - English dictionary for spelling corrections, definitions
 - Compilers: symbol tables (list of names and related info)
 - Network routers: port number -> socket id
 - virtual memory: virtual address -> physical address

Conclusions

- Hashing is an efficient way to keep average cost of operations to $O(1)$
- Collisions are unavoidable in practice and are solved by chaining
- Worse case $\Theta(n)$
- We have some simple ways to construct “good” hash functions
- Hashes are not first choice if worst case behaviour is important
 - Better worst case: Van Emde Boas Trees $O(\log \log n)$ [not covered here]
- “Hashes kill caches”